

Data Science for Economists

Introduction to Machine Learning

Mox Ballo

University of the Philippines

April 23, 2025



SCHOOL OF ECONOMICS
UNIVERSITY OF THE PHILIPPINES
DILIMAN, QUEZON CITY

① Introduction to Machine Learning

② Components of Learning

③ Types of Learning

④ Learning Feasibility

⑤ Error and Noise

1 Introduction to Machine Learning

2 Components of Learning

3 Types of Learning

4 Learning Feasibility

5 Error and Noise

Overview

- Learning = improving performance at a task through experience
- Formal set-up: input x , output y , unknown target function f
- Goal: learn an approximation $g(x)$ using data
- Example: Forecasting electricity demand from past weather and consumption

Set-up of the Learning Problem

- **Inputs:** $x \in \mathcal{X}$ — observed variables (e.g., hour, weather, previous demand)
- **Outputs:** $y \in \mathcal{Y}$ — quantity to predict (e.g., next-hour electricity load)
- **Input Space:** \mathcal{X} — all possible combinations of input features
- **Output Space:** \mathcal{Y} — all valid values the output can take
- **Unknown Target Function:** $f : \mathcal{X} \rightarrow \mathcal{Y}$ — true, hidden mapping
- **Learning Objective:** Find an approximation $g \in \mathcal{H}$ such that $g(x) \approx f(x)$

Learning Problems in Economics

- **Demand forecasting:** Predict electricity load based on time, weather, and historical usage
- **Price elasticity estimation:** Learn how quantity demanded responds to price changes using observed market data
- **Consumer choice modeling:** Predict product choice given consumer attributes and product features

- 1 Introduction to Machine Learning
- 2 Components of Learning
- 3 Types of Learning
- 4 Learning Feasibility
- 5 Error and Noise

Input Space and Feature Selection

Problem: Predict hourly electricity demand to optimize production.

- **Vector of Inputs** (x_n): hour of day, day of week, temperature, humidity, historical consumption
- **Input Space** (\mathcal{X}):

$$\mathcal{X} = \{0-23\} \times \{\text{Mon-Sun}\} \times \mathbb{R}^+ \times [0, 100] \times \mathbb{R}^+$$

Output and Hypothesis Space

- **Output (y_n):** The observed electricity demand, e.g. in megawatts (MW), at time n .
- **Output Space (\mathcal{Y}):** Since demand cannot be negative, we model outputs in the space of non-negative real numbers: $\mathcal{Y} = \mathbb{R}^+$.
- **Hypothesis Space (\mathcal{H}):** The set of functions $g : \mathcal{X} \rightarrow \mathcal{Y}$ that map input features (e.g., time of day, weather) to a predicted demand value.
 - Linear models: $\mathcal{H} = \{g(x) = w^\top x\}$
 - Polynomial regression: \mathcal{H} includes higher-order terms to capture nonlinear trends.
 - Nonparametric models: GAMs or decision trees allow more flexible hypothesis spaces.

Learning Algorithm

- **Goal:** Use training data $\mathcal{D} = \{(x_n, y_n)\}_{n=1}^N$ to choose a hypothesis $g \in \mathcal{H}$ that approximates the unknown target function f .
- **Learning Algorithm (\mathcal{A}):** A systematic procedure that selects the final hypothesis:

$$g = \mathcal{A}(\mathcal{D}) \in \mathcal{H}$$

- **Examples for electricity demand:**
 - Linear regression (OLS, Ridge, Lasso)
 - Tree-based models (Random Forest, XGBoost)
 - Neural networks for time series
- **Output:** A learned function g that generalizes well to new inputs x_{new}

Error Measures

- **Final Hypothesis:** The learned function $g(x) = f(x; \hat{\beta})$ aims to approximate the unknown target f .
- **In-Sample Error (Training Error):**

$$E_{\text{in}}(g) = \frac{1}{N} \sum_{n=1}^N (g(x_n) - y_n)^2$$

Measures how well g fits the training data.

- **Out-of-Sample Error (Generalization Error):**

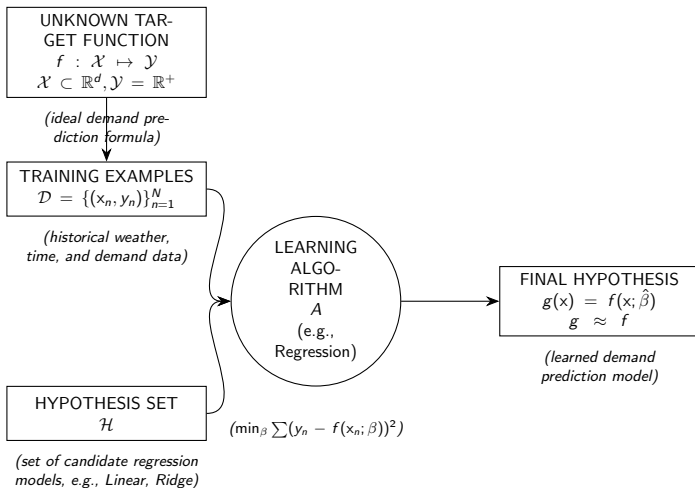
$$E_{\text{out}}(g) = \mathbb{E}_{x,y} \left[(g(x) - y)^2 \right]$$

Measures how well g performs on new, unseen data.

Error Measures

- **Goal of Learning:** Minimize $E_{\text{out}}(g)$, not just $E_{\text{in}}(g)$. This ensures good generalization.
- **Common Metrics:**
 - **MSE (Mean Squared Error)** — most common for regression.
 - **MAE (Mean Absolute Error)** — less sensitive to outliers.
 - **MAPE (Mean Absolute Percentage Error)** — often used in electricity demand forecasting.
- **Why it matters:** A low in-sample error doesn't guarantee a low out-of-sample error — hence the need for validation techniques (e.g. cross-validation).

Basic Set-up: Electricity Demand Prediction



A Simple Learning Model

- Consider forecasting the electricity demand for a given hour in a day, given weather and time-based conditions.
- The training examples $\mathcal{D} = \{(x_n, y_n)\}_{n=1}^N$ come from historical demand logs.
- We choose a hypothesis set \mathcal{H} , and a learning algorithm \mathcal{A} , to find a function $g \approx f$ that maps inputs x to predicted demand y .
- A simple model: assume demand can be predicted as a weighted sum of features:

$$g(x) = \sum_{i=1}^d w_i x_i + b$$

- for simplicity, instead of β , we will use w

What's in the Input Vector x ?

- $x \in \mathbb{R}^d$, where each component of x represents:
 - x_1 : Hour of the day (0–23)
 - x_2 : Day of the week
 - x_3 : Temperature
 - x_4 : Humidity
 - x_5 : Lagged demand (e.g. demand one hour ago)
- Output $y \in \mathbb{R}^+$: actual electricity demand in megawatts (MW)
- Our goal: learn a function $g(x)$ that gives accurate predictions of y

Linear Hypothesis for Demand

- We define our hypothesis space \mathcal{H} to consist of linear models:

$$h(\mathbf{x}) = \sum_{i=1}^d w_i x_i + b$$

- w_i : learned importance of each input feature (e.g., temperature, time)
- b : bias term, reflects baseline electricity demand
- Compact form:

$$h(\mathbf{x}) = \mathbf{w}^\top \mathbf{x} + b$$

- After training, we get final hypothesis:

$$g(\mathbf{x}) = h(\mathbf{x}; \hat{\mathbf{w}})$$

Learning the Parameters: ERM

- We use the training data to estimate weights \hat{w} and bias \hat{b}
- Objective: minimize squared error (Empirical Risk Minimization)

$$\min_{w,b} \frac{1}{N} \sum_{n=1}^N \left(y_n - (w^\top x_n + b) \right)^2$$

- This selects the best-fitting line (or hyperplane) in the training data
- The result is our final hypothesis:

$$g(x) = w^\top x + b$$

Learning vs. Design

- In econometrics, models are often specified from theory: *design-based*.
- In machine learning, we let the data suggest the best hypothesis: *data-driven*.
- Tradeoff:
 - Design-based: interpretable, theory-aligned, but may misrepresent data.
 - Learning-based: better predictive accuracy, but harder to interpret and justify causally.
- Our goal is to understand when and how data-driven models (like this) can augment economic analysis.


Digression: Perceptrons

- A perceptron is a simple binary classifier.
- It takes multiple inputs (features), computes a weighted sum, and applies a decision rule:

$$h(x) = \text{sign}(w^T x + b)$$

- Think of it as a “yes/no” decision boundary — like deciding if the grid is stable or at risk.
- Inspired by how neurons in the brain work: inputs fire, and if the signal is strong enough, there's an activation (output = 1).

Illustration: 2-Feature Case



Digression Summary

- A perceptron is one of the simplest machine learning models.
- It separates the space into two regions using a line (or hyperplane in higher dimensions).
- It's fast and interpretable — great for simple decisions like:
 - Is the power grid at risk of overload?
 - Should we activate a reserve power plant?
- But it only works well if the data can be separated by a straight line (linearly separable).

Perceptron-Style Binary Classification

- In some energy markets, the goal is to make a **binary decision**:

$$y \in \{-1, +1\} \quad (\text{e.g., brownout risk?})$$

- Input vector x might include:
 - x_1 : Forecasted demand (MW)
 - x_2 : Available generation capacity
 - x_3 : Grid frequency deviation
 - x_4 : Weather forecast (e.g., typhoon alert)
- Output: $y = +1$ means “risk of brownout” and $y = -1$ means “normal operation”

Perceptron Model

- Define hypothesis:

$$h(x) = \text{sign}(w^T x + b)$$

- Interpretation:
 - If $h(x) = +1$: trigger preventive measures (e.g., activate reserve plants)
 - If $h(x) = -1$: system is stable
- The learning algorithm finds weights w that separate high-risk vs. low-risk conditions.
- If input space is linearly separable, the Perceptron algorithm converges to a solution.

Philippine Examples of Binary Energy Decisions

- **Scenario 1: Load Shedding**
 - NGCP needs to decide whether to issue yellow or red alerts
 - Predict based on forecast demand and supply margin
- **Scenario 2: Rotational Brownouts**
 - During the dry season, hydroelectric capacity is low
 - Predict risk of brownout in provinces like Palawan or Mindoro
- **Scenario 3: Distributed Energy Dispatch**
 - If risk $h(x) = +1$, then dispatch diesel backup generators in off-grid areas

① Introduction to Machine Learning

② Components of Learning

③ Types of Learning

④ Learning Feasibility

⑤ Error and Noise

Supervised Learning

- In supervised learning, the algorithm is given input-output pairs (x, y) .
- The goal is to learn a function $g(x)$ that predicts y from x .
- Two major types:
 - **Regression:** Predict continuous output (e.g. electricity demand in MW)
 - **Classification:** Predict binary or categorical output (e.g. default vs. no default)
- The algorithm is “supervised” by the correct answer.
- **Examples:**
 - Predicting household electricity consumption based on time and weather
 - Forecasting inflation using macroeconomic indicators
 - Predicting approval for microfinance loans based on applicant data

Active Learning in Supervised Learning)

- In standard supervised learning, the algorithm passively receives labeled data.
- In **active learning**, the learner can query for labels — it chooses which x to ask for a y .
- Motivation: labels may be expensive or time-consuming to obtain.
- **Example:** Smart meters collect hourly consumption (x), but manual surveys provide income class (y).
- Active learning asks: "Which households should we label (survey) to improve our model fastest?"
- **Examples:**
 - Reducing survey costs in poverty targeting
 - Sampling households for tariff adjustment studies
 - Choosing which firms to audit for competition enforcement

Online Learning

- In **online learning**, the algorithm updates its model one observation at a time.
- Useful for streaming data or when retraining on full data is costly.
- At each step t , the learner:
 - ① Gets input x_t
 - ② Predicts $\hat{y}_t = g_t(x_t)$
 - ③ Receives actual y_t and incurs loss
 - ④ Updates model to g_{t+1}
- **Examples:**
 - Real-time pricing forecasts in electricity markets (e.g., WESM)
 - Updating inflation predictions with streaming macroeconomic indicators
 - Learning user behavior in mobile payment platforms

Unsupervised Learning

- In unsupervised learning, we are given only inputs x_1, x_2, \dots, x_N , with no labels or targets.
- objective: discover patterns, structures, or relationships within the data.
- The algorithm is “unsupervised” — it must learn without knowing the correct answer.
- Useful when labels are unavailable or ambiguous.
- **Main types:**
 - **Clustering** (e.g. k-means, hierarchical clustering)
 - **Dimensionality reduction** (e.g. PCA, t-SNE)

Clustering

- Clustering groups similar observations together based on a distance or similarity metric.
- Examples:
 - K-means clustering
 - Hierarchical clustering
 - DBSCAN
- Clustering reveals hidden groupings in the data — no labels are needed.
- **Examples:**
 - Segmenting consumers based on electricity usage patterns
 - Grouping municipalities by poverty and infrastructure indicators
 - Identifying similar firms in competition analysis

Dimensionality Reduction

- Reduces the number of features while retaining the most important variation in the data.
- Techniques:
 - Principal Component Analysis (PCA)
 - t-SNE for visualization
 - Autoencoders (neural networks)
- Helps:
 - Visualize high-dimensional data
 - Remove noise or redundancy
 - Improve downstream supervised learning
- **Example:**
 - Reducing 50 household expenditure categories into 2–3 latent "lifestyle" dimensions

Supervised vs. Unsupervised Learning

Aspect	Supervised Learning	Unsupervised Learning
Input	Features + labels (x, y)	Features only (x)
Goal	Predict labels	Find structure
Feedback	Known correct outputs	No direct feedback
Methods	Regression, classification	Clustering, PCA

Reinforcement Learning

- In reinforcement learning, an agent learns by interacting with an environment.
- It receives feedback in the form of rewards and learns a policy to maximize long-term reward.
- No fixed training set — learning happens over time via trial and error.
- **Examples:**
 - Dynamic pricing of electricity under real-time market conditions
 - Adaptive congestion pricing in transport networks
 - Policy simulation: learning optimal subsidy schemes under changing economic conditions
- Especially useful for problems involving time, strategy, and sequential decisions.

Key Components of Reinforcement Learning

- **Agent:** Learner or decision-maker (e.g., National Grid operator)
- **Environment:** World with which the agent interacts (e.g., electricity grid)
- **State (s):** Current situation observed by the agent (e.g., demand level, reserve margin)
- **Action (a):** Decision the agent can take (e.g., dispatch reserve, adjust price)
- **Reward (r):** Feedback signal after taking action (e.g., blackout avoided = +10; blackout = -100)
- **Policy (π):** Mapping from state to action: $a = \pi(s)$
- **Goal:** Learn the best policy π^* that maximizes long-term cumulative reward

What is the “Training Set” in RL?

- Unlike supervised learning, RL does not receive labeled pairs (x, y) .
- Instead, the agent observes a sequence of experience tuples:

$$(s_t, a_t, r_t, s_{t+1})$$

- This tells the agent:
 - “I was in state s_t ”
 - “I chose action a_t ”
 - “I received reward r_t ”
 - “The new state is s_{t+1} ”
- Learning happens by updating value estimates or policies from this feedback.

Where is RL Useful?

- Reinforcement Learning is ideal when:
 - The agent must make sequential decisions
 - The environment changes dynamically
 - Outcomes depend on both present and past decisions
- **Possible Applications:**
 - Energy markets: dispatch decisions, real-time pricing
 - Transport economics: dynamic tolling, congestion response
 - Social policy: testing subsidy levels under uncertain behaviors
 - Public health: optimizing allocation of limited vaccines over time
- RL is often used in simulation-based environments when real-world experimentation is costly.

➊ Introduction to Machine Learning

➋ Components of Learning

➌ Types of Learning

➍ Learning Feasibility

➎ Error and Noise

Is Learning Always Feasible?

- Learning is only useful if we can generalize from the data.
- Key question: If a model works well on training data, will it also work on new, unseen data?
- This is the difference between:
 - **Memorization** — fitting the training data exactly
 - **Generalization** — capturing the true underlying pattern
- Learning is feasible when we can ensure small error on unseen data (out-of-sample).
- This requires assumptions on:
 - How the data is generated
 - The complexity of the hypothesis space

Condition 1: How the Data is Generated

- We assume the training data $(x_1, y_1), \dots, (x_N, y_N)$ are drawn independently from the same probability distribution \mathcal{D} .
- This is the i.i.d. assumption (independent and identically distributed).
- Why it matters:
 - If data changes over time (non-stationary), past experience may not generalize.
 - If observations are correlated (e.g., clustered or networked), standard learning guarantees break down.
- **Examples:**
 - Forecasting electricity demand works if usage patterns are stable — not during blackouts or crises.
 - Learning a pricing model for MC Taxi fares only works if commuter behavior is consistent.
- Takeaway: For learning to generalize, the future must “look like” the past.

Condition 2: Complexity of the Hypothesis Space

- A hypothesis space \mathcal{H} is the set of functions we allow the learner to choose from.
- If \mathcal{H} is too rich (too many flexible models), it can perfectly fit the training data but fail to generalize.
- The more complex the hypothesis set, the more data we need to ensure learning is feasible.
- Complexity can be measured by:
 - VC Dimension
 - Number of parameters
 - Model structure (e.g., depth of trees, layers of a neural net)
- **Examples:**
 - A flexible forecasting model with dozens of lag terms, interaction effects, and splines may overfit if sample size is small.
 - Simpler models (e.g., linear regression with 3–4 predictors) are more robust in small datasets.

• **Trade-off:** More expressive models vs. risk of overfitting.

Why Probability Matters in Learning

- Learning is about drawing conclusions from finite samples.
- Even if we fit the training data well, we might just be overfitting — capturing noise.
- We use probability to:
 - Estimate the likelihood of making a mistake on unseen data
 - Bound the difference between in-sample error and out-of-sample error
- **Example:** If we train a demand forecasting model on dry-season data, can it generalize to rainy season?
- Formal tools: concentration inequalities, generalization bounds

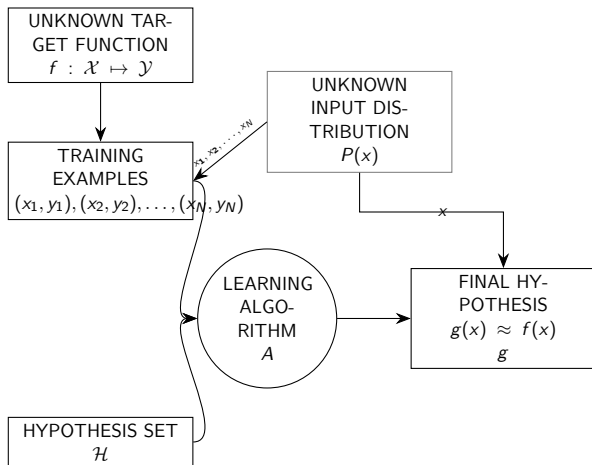
Putting It All Together

- Learning is feasible when:
 - ① The data is i.i.d. from a fixed distribution
 - ② The hypothesis space has controlled complexity
- We can then bound the difference between:

$$E_{\text{out}}(g) \approx E_{\text{in}}(g)$$

- This is the heart of generalization.
- In economics and policy:
 - Make sure your training data is relevant to your policy domain.
 - Don't fit overly flexible models without enough observations.

Learning Setup with Probability



The VC Dimension

- The **VC Dimension** (Vapnik–Chervonenkis) measures the complexity of a hypothesis set.
- Informally: it's the largest number of points that can be “shattered” (i.e., labeled in all possible ways).
- **Higher VC dimension** → more expressive models, but also higher risk of overfitting.
- **Lower VC dimension** → simpler models, may underfit.
- Goal: find a balance where the hypothesis set is expressive enough but still generalizes well.

Examples of VC Dimension

- **Threshold function on the real line:** VC dimension = 1
- **Linear classifiers in 2D:** VC dimension = 3
- **Perceptron with d inputs:** VC dimension = $d + 1$
- **Implication for electricity demand prediction:**
 - If you use too many features (hour, temperature, weather alerts, device types, etc.), you may overfit.
 - VC dimension helps guide the size of the hypothesis space.

Summary: When is Learning Feasible?

- Learning is feasible when:
 - The hypothesis space is not too large
 - The number of training examples is sufficient
 - The data distribution is stable over time
- Probability theory helps us quantify generalization error.
- VC dimension tells us when a model is too complex.
- **In economics:** Learning is feasible when:
 - You don't have too many instruments relative to observations
 - Consumer behavior is relatively stable

① Introduction to Machine Learning

② Components of Learning

③ Types of Learning

④ Learning Feasibility

⑤ Error and Noise

Understanding Error in Learning

- In supervised learning, we want our hypothesis g to approximate the true function f .
- But even the best model g will make some mistakes.
- We distinguish between:
 - **In-sample error** E_{in} : How well g performs on training data
 - **Out-of-sample error** E_{out} : How well g performs on new, unseen data
- Goal: minimize E_{out} , not just memorize training data.

Common Error Measures

- **Mean Squared Error (MSE):**

$$\text{MSE} = \frac{1}{N} \sum_{n=1}^N (g(x_n) - y_n)^2$$

- **Mean Absolute Error (MAE):**

$$\text{MAE} = \frac{1}{N} \sum_{n=1}^N |g(x_n) - y_n|$$

- **Binary Classification Error:**

$$\text{Error Rate} = \frac{1}{N} \sum_{n=1}^N 1 [g(x_n) \neq y_n]$$

- Choice depends on the task and whether large errors should be penalized more.

What is Noise?

- **Noise** refers to the randomness or uncertainty in the relationship between x and y .
- Two data points with identical x may have different y values.

$$y = f(x) + \text{noise}$$

- Noise is often due to:
 - Measurement errors
 - Omitted variables
 - Human unpredictability
- Noise sets a lower bound on the best achievable performance.
- Even the best model can't predict noise — only patterns.

Noisy Targets in Economics

- Many economic outcomes are inherently noisy:
 - Electricity consumption depends on mood, appliance use, etc.
 - Household income may fluctuate due to informal jobs
 - Self-reported survey data often contains inaccuracies
- As a result:
 - Don't expect perfect accuracy — measure error realistically
 - Focus on robustness and generalization, not fitting every detail
- **Design implication:** Simpler models may outperform complex ones in noisy settings.

Summary:

- Total error can be broken down into:

$$\text{Total Error} = \text{Bias}^2 + \text{Variance} + \text{Noise}$$

- **Bias:** Error due to wrong assumptions (underfitting)
- **Variance:** Error due to model sensitivity to data (overfitting)
- **Noise:** Irreducible error from randomness in the data
- As economists and data scientists, our job is to minimize bias and variance — not noise.

Bias-Variance Tradeoff

- When we train a model, the goal is to predict well — not just memorize.
- But prediction error has **three ingredients**:

$$\text{Total Error} = \underbrace{\text{Bias}^2}_{\text{wrong model}} + \underbrace{\text{Variance}}_{\text{too flexible}} + \underbrace{\text{Noise}}_{\text{randomness}}$$

- **Bias** is the error from using a model that's too simple (e.g., a straight line for something curvy).
- **Variance** is the error from using a model that's too complex — it overreacts to small changes in the data.
- **Noise** is unavoidable randomness (e.g., people leaving the aircon on by accident).
- **Tradeoff:**
 - A simple model (low variance, high bias) may miss real trends.
 - A complex model (low bias, high variance) may chase noise.
- We want a balance — not too simple, not too wiggly.

Visualizing Bias vs. Variance

Low Bias,
Low Variance



Accurate & consistent

High Bias,
Low Variance




Inaccurate, consistent

Low Bias,
High Variance



Accurate (avg),
inconsistent

High Bias,
High Variance



Inaccurate & inconsistent